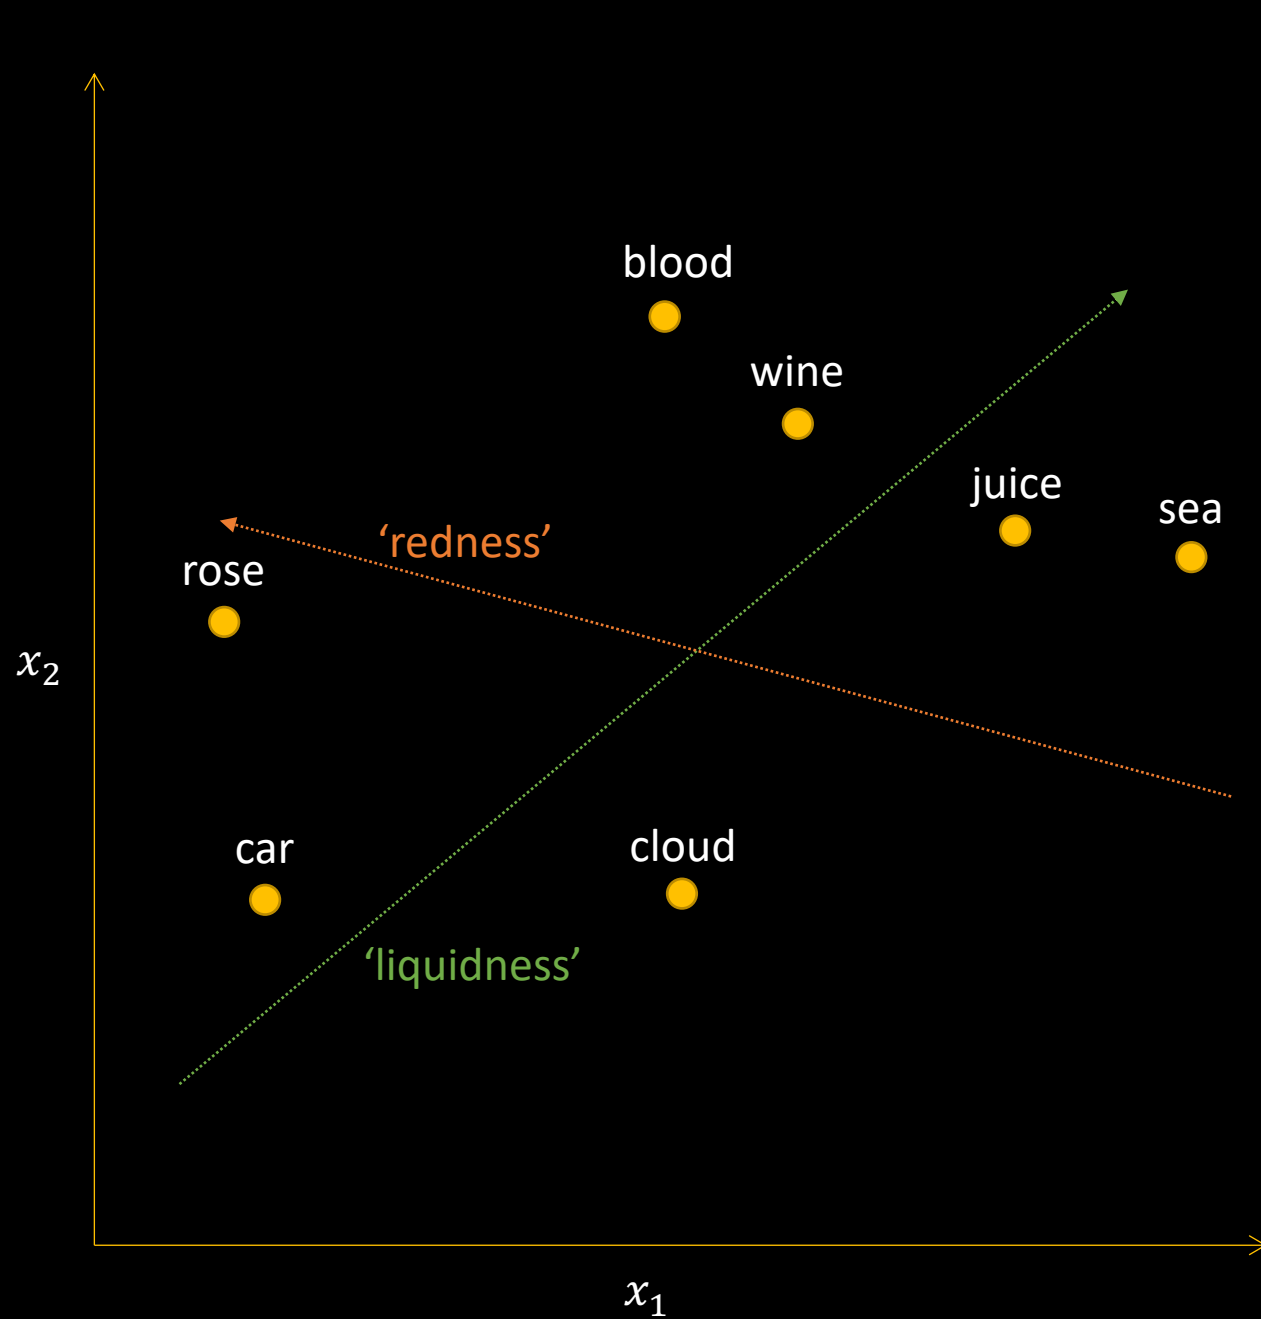


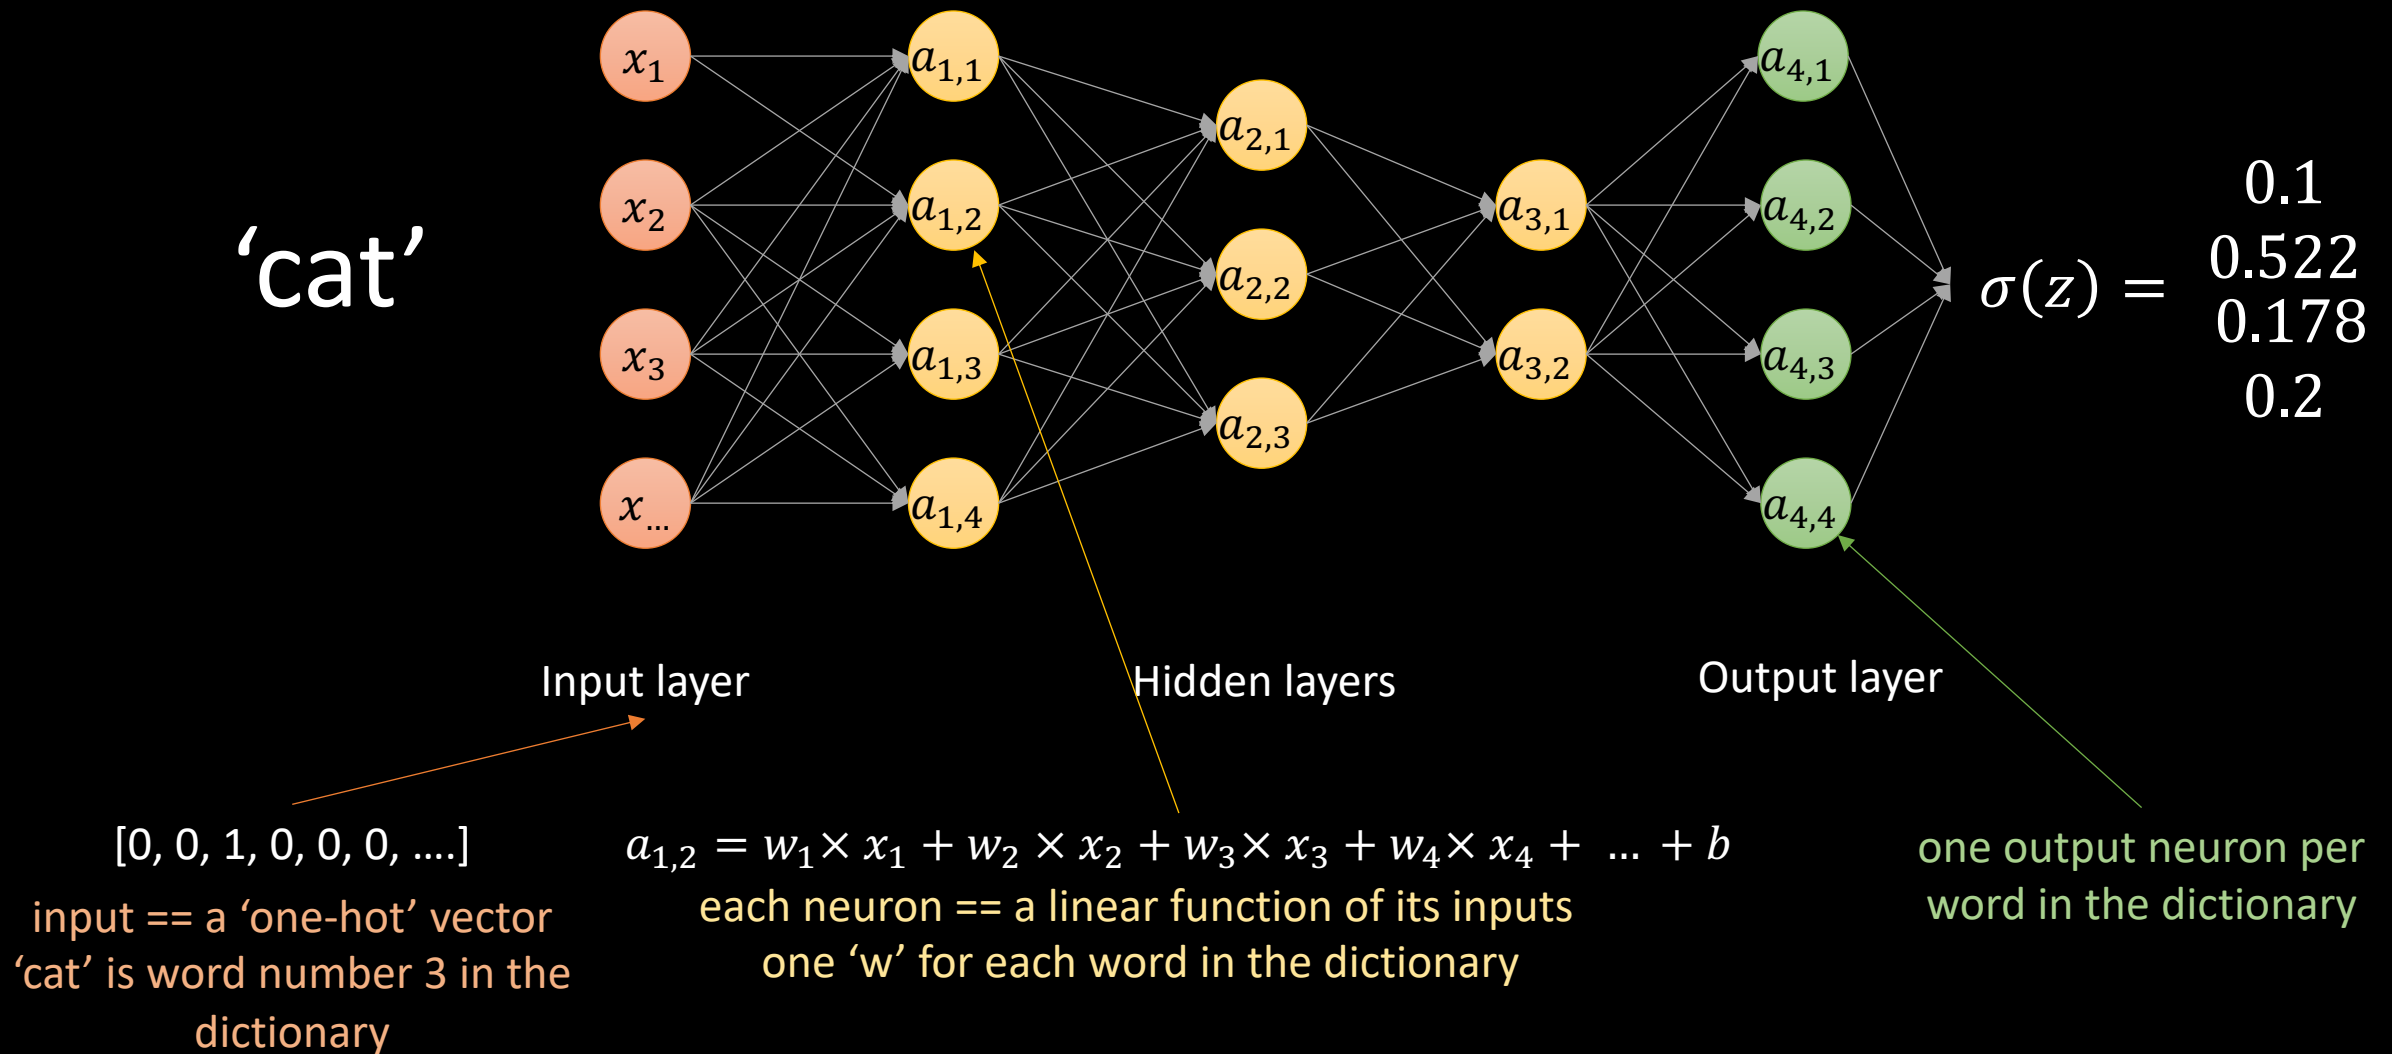
# Word Vectors




$$rose = \begin{pmatrix} 1.2 \\ 6.9 \end{pmatrix} = \mathbb{R}_2$$

<i>word</i>	$x_1$	$x_2$
rose	1.2	6.9
blood	5.1	9.3
car	1.5	3.9
sea	9.3	7.1
wine	6.0	8.5
juice	7.9	7.2
cloud	4.9	4.0

# Can we predict the next word in a sentence?



'cat' == word 4

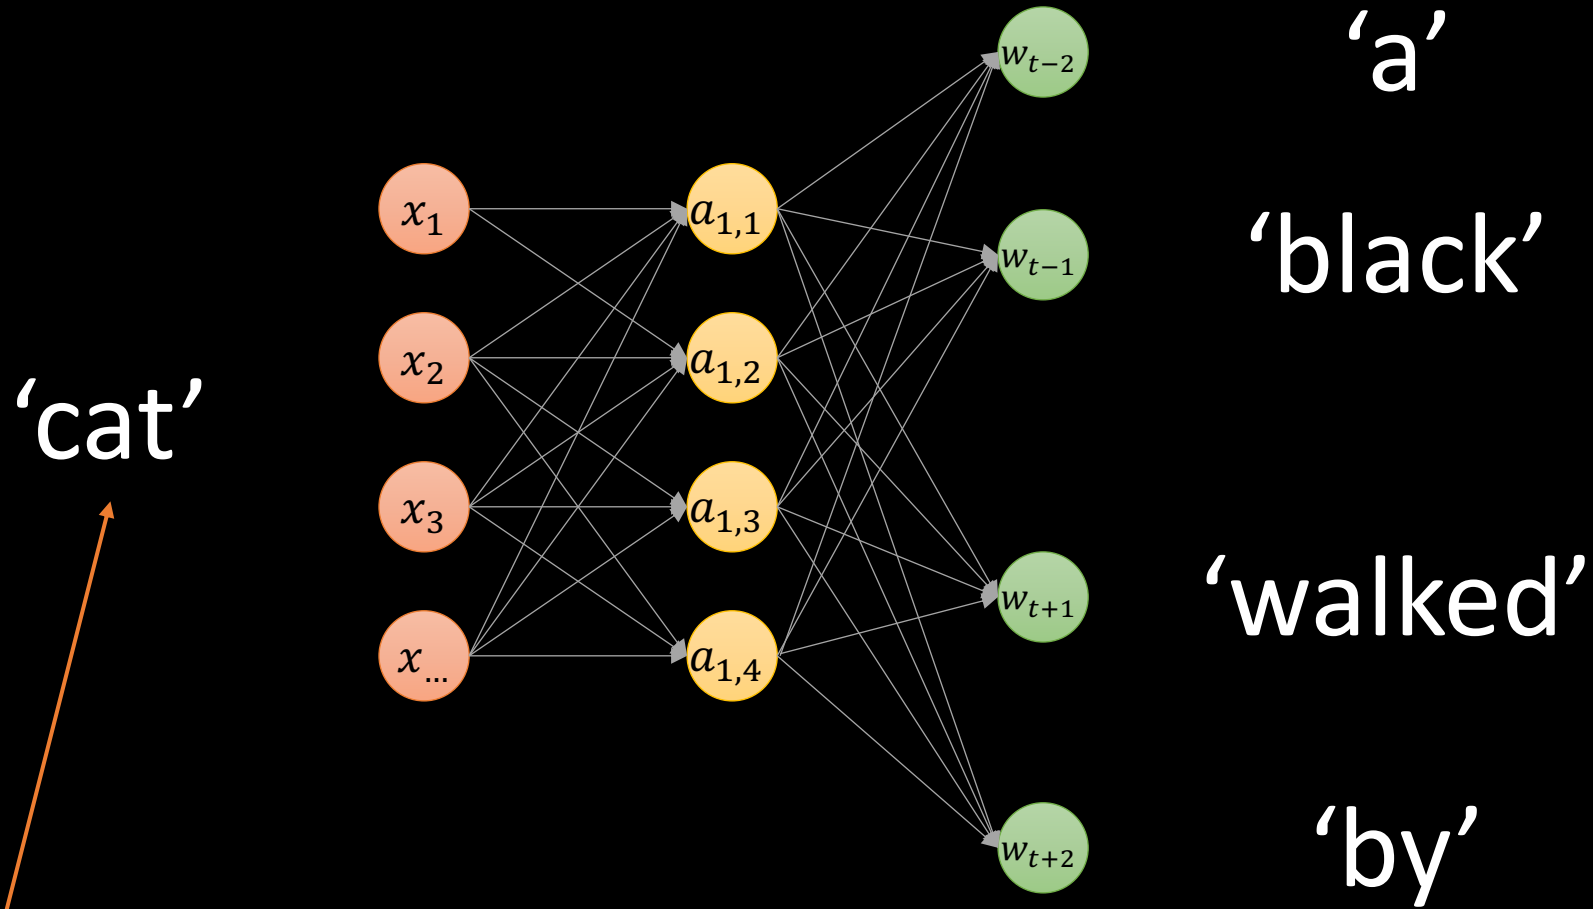


This set of numbers represents the network's idea of the word 'cat'.

$$a_{1,2} = w_1 \times x_1 + w_2 \times x_2 + w_3 \times x_3 + w_4 \times x_4 + \dots + b$$

each neuron == a linear function of its inputs  
one 'w' for each word in the dictionary

# The 'word2vec' skip-gram model

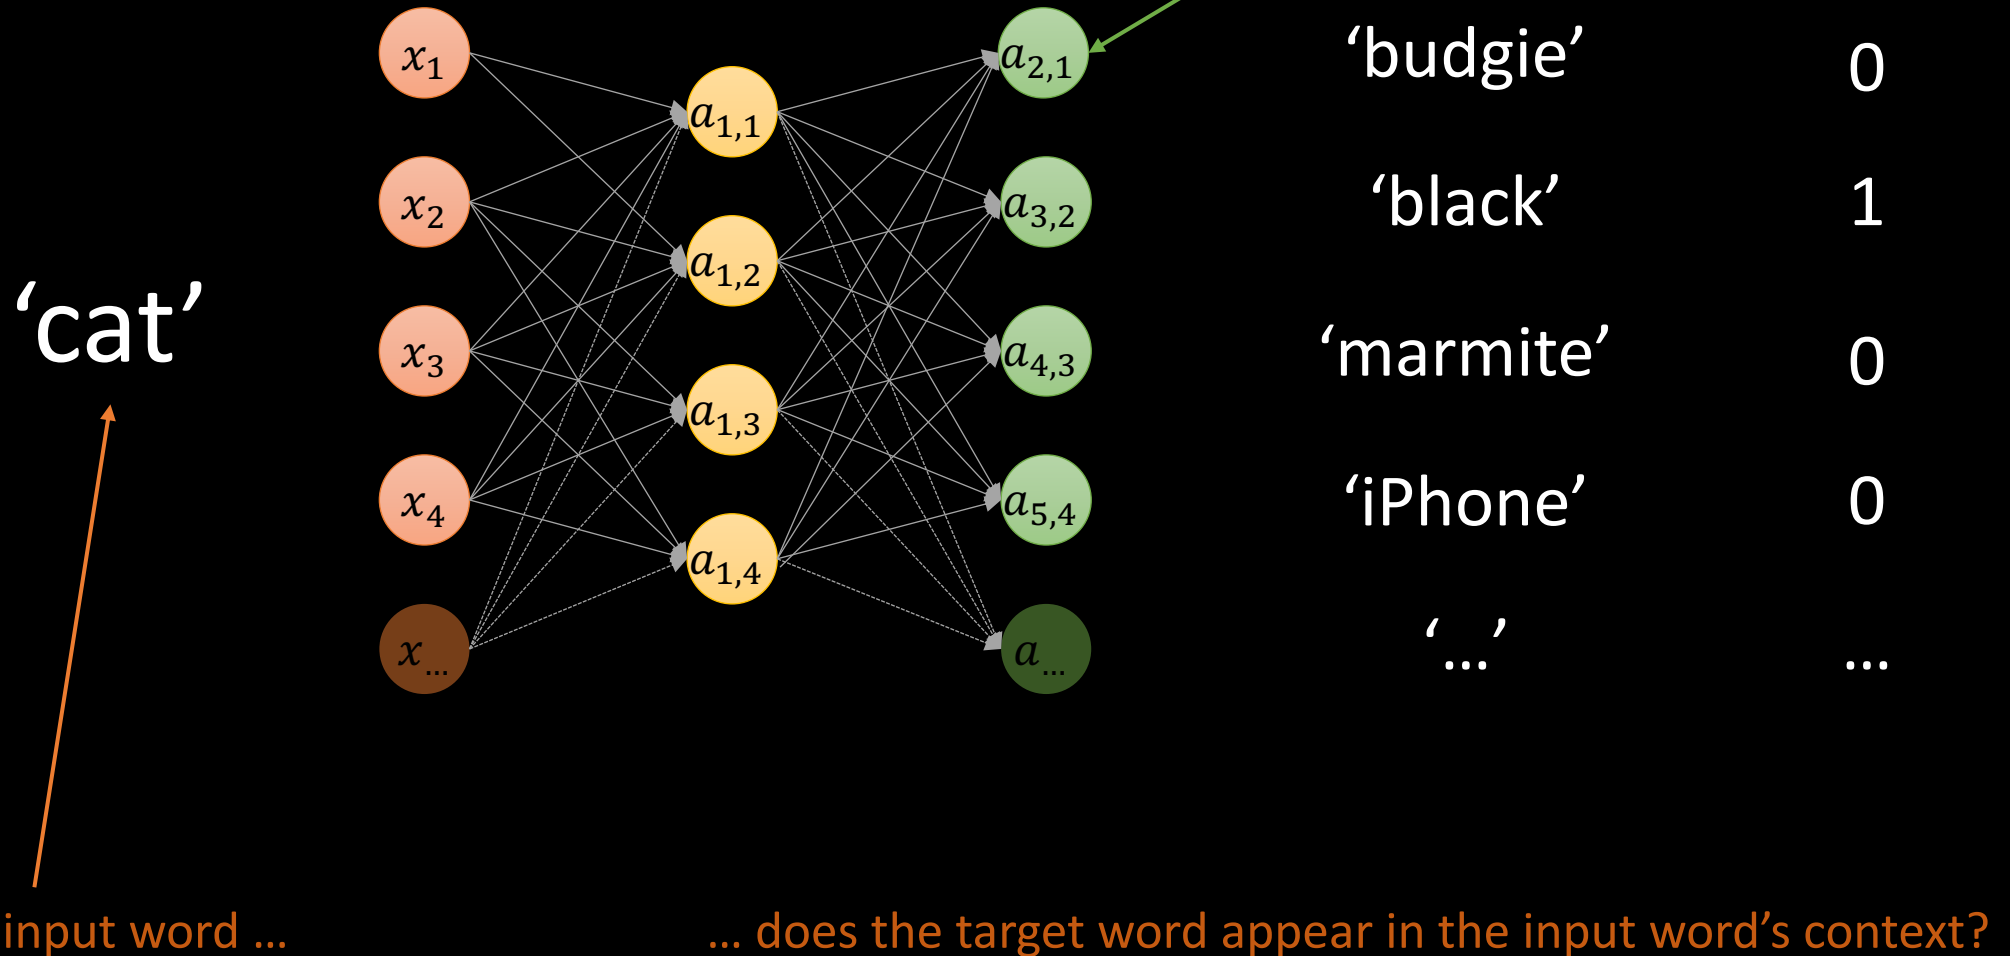


Given this word ...

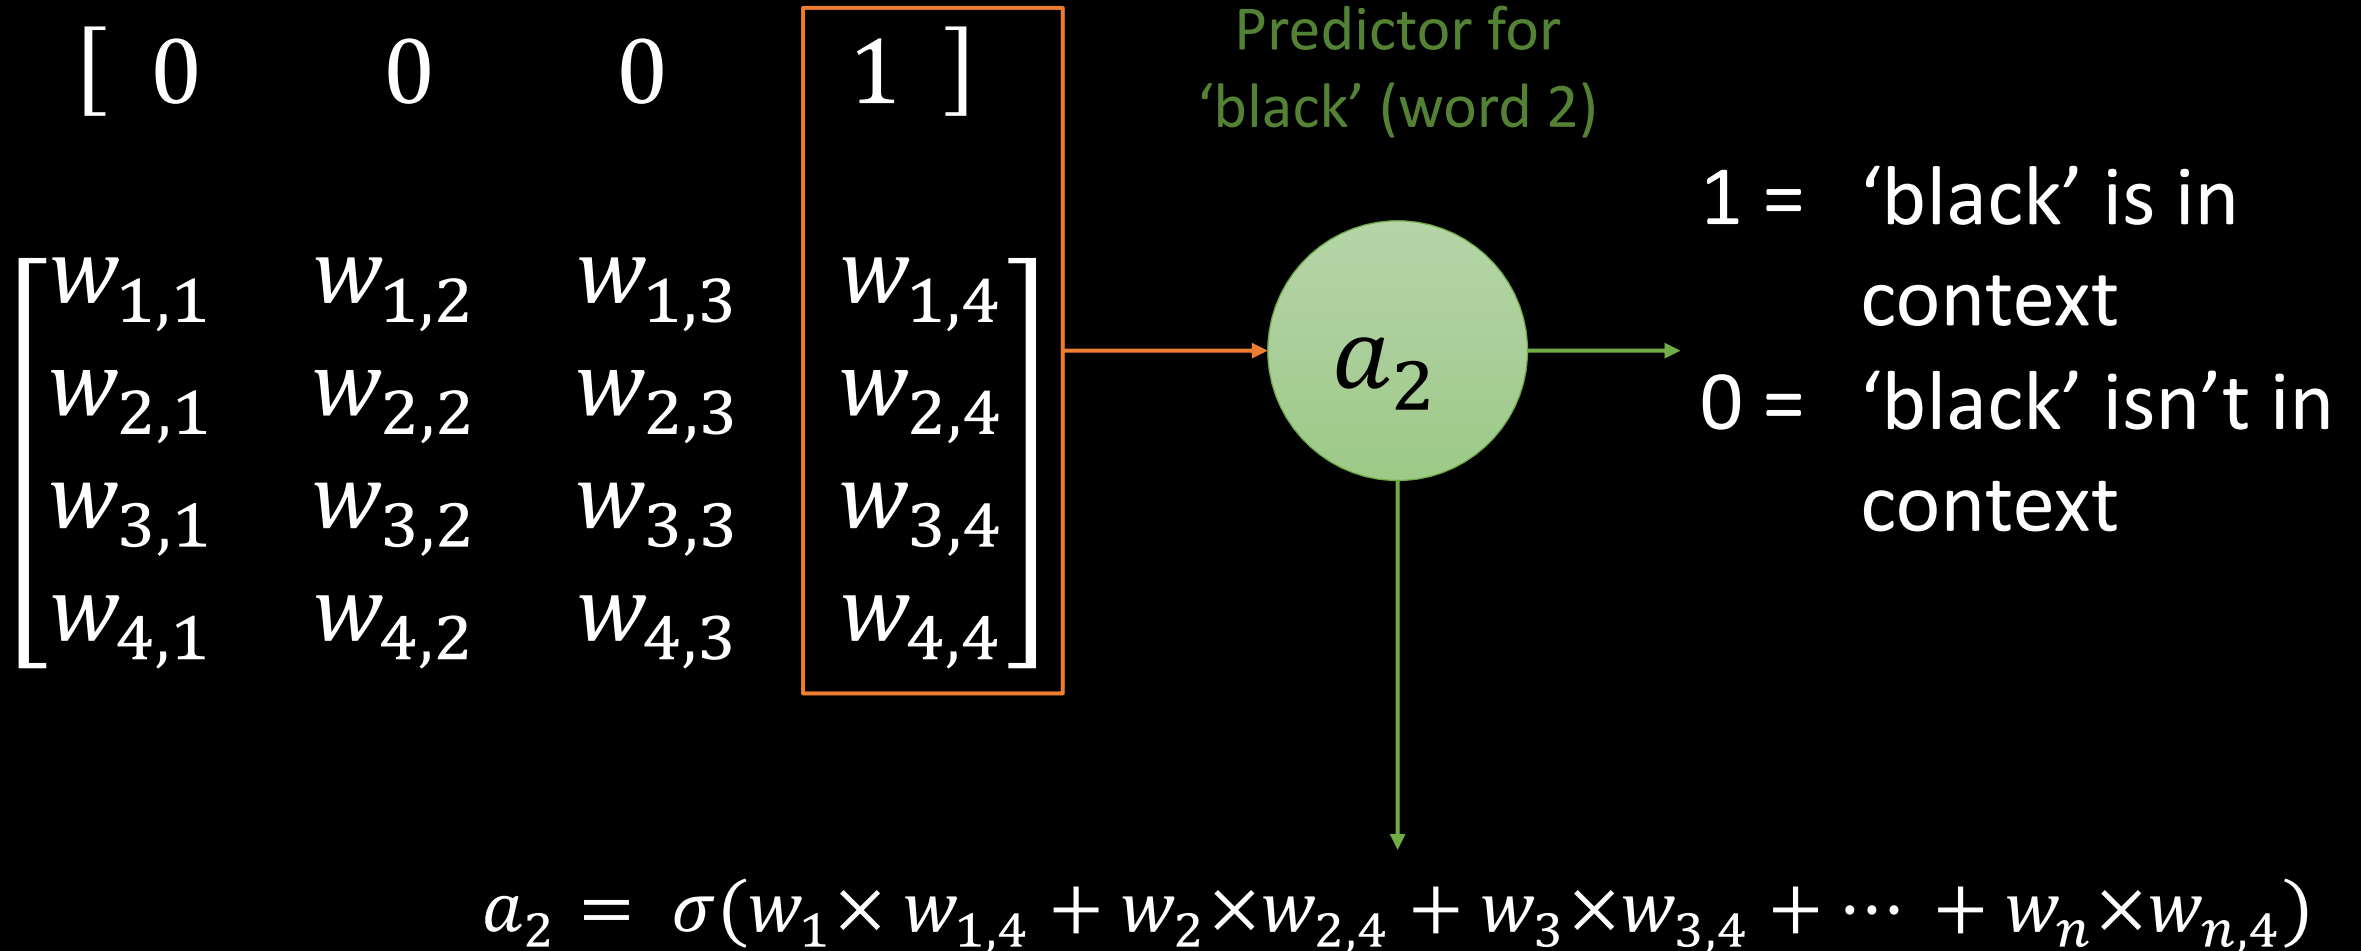
... predict which words appear nearby (window\_size == 2)

# 'word2vec' with negative sampling

One prediction neuron for each word in vocabulary  
(10,000 words = 10,000 neurons)



'cat' == word 4



# Training data for negative sampling

‘When I looked **out**, a black **cat** crossed the road, and purred.’

window\_size == 3


input word	target word	value
cat	out	1
cat	blue	0
cat	spandex	0
cat	ballet	0
cat	thingamy	0

Actual word

Randomly sampled  
negative examples



What can  
you do with  
word  
vectors?



Tomas Mikolov and others, 'Distributed Representations of Words and Phrases and Their Compositionality', in *Advances in Neural Information Processing Systems 26*, ed. by C. J. C. Burges and others (Curran Associates, Inc., 2013), pp. 3111–19.